

1. How might resource constraints (like RAM, CPU/GPU availability) impact the performance and usability of locally deployed LLMs?

LLMs demand substantial computational and memory resources to function effectively. They require significant RAM to load model weights and process input-output tensors during inference. When system memory is insufficient, the model may experience drastically slower inference times and fail to load altogether. Similarly, disk speed influences performance, slow hard drives increase model loading times compared to solid-state drives, making real-time applications less feasible. Additionally, LLMs occupy large storage spaces, often exceeding 10–20 GB for mid-sized models like a 7B parameter system in half-precision. These factors make deployment on consumer-grade devices technically challenging and limit usability for applications requiring continuous responsiveness.

GPU availability is another critical determinant of performance. GPUs are optimized for parallel matrix operations essential to transformer-based models, offering massive speedups over CPUs. Without a capable GPU, inference and fine-tuning rely solely on CPUs, which are often 10–100 times slower depending on the task. Limited CPU cores also constrain parallel token generation, reducing throughput, for example, a standard consumer CPU may generate only 1–2 tokens per second, whereas a high-end GPU can exceed 50 tokens per second. Consequently, users facing hardware constraints may need to opt for smaller models (e.g., 3B instead of 70B parameters), sacrificing fluency, reasoning ability, and contextual understanding. These limitations make locally deployed LLMs impractical for real-time applications such as chatbots.

2. What ethical responsibilities do developers have when deploying and customizing local LLMs?

Developers have a duty to ensure fairness, transparency, and accountability when deploying or customizing local LLMs. They must audit training and fine-tuning data to identify and reduce biases, especially those involving gender, race and socioeconomic representation. Techniques such as dataset balancing can help mitigate these issues and prevent models from reinforcing stereotypes. Continuous evaluation is also essential to detect and correct problematic behaviors as the model interacts with real users.

Equally important is maintaining transparency and user trust. Developers should clearly document model architectures, data sources, and modification processes to allow others to assess limitations and potential risks. Protecting user privacy is another ethical priority, ensuring that local inference or logging does not expose sensitive information. Neglecting these responsibilities can lead to misuse, bias propagation, and erosion of public confidence in AI systems. Ethical stewardship means designing local LLMs that prioritize human well-being as much as technical performance.

3. Discuss the potential benefits and limitations of using AI-assisted tools like Roo Code for data analysis tasks, and what aspects of this process feels different from conventional coding projects.

AI-assisted tools such as Roo Code enhance data analysis by automating routine tasks, suggesting efficient code, and identifying errors in real time. They help users translate natural language instructions into executable code, accelerating workflow and improving accessibility for beginners. For experienced analysts, these tools can boost productivity, support rapid prototyping, and maintain consistency in coding practices. The interactive, conversational interface also fosters a more intuitive and exploratory approach to programming compared to traditional manual coding.

Unlike conventional coding, which emphasizes deliberate planning and debugging, AI-assisted workflows rely more on collaboration with the model, demanding human oversight to verify accuracy and maintain methodological rigor. Thus, while Roo Code enhances efficiency, it must be used critically to ensure analytical validity and reliability.

4. How does virtualisation contribute to the efficiency of cloud computing?

Virtualization allows multiple virtual machines to run on a single physical server by abstracting hardware resources. This enables efficient utilization of CPU, memory, and storage by allocating them dynamically to different users and applications as needed. Cloud providers can scale resources up or down seamlessly, reducing idle capacity and lowering operational costs. It also simplifies deployment and maintenance, allowing rapid provisioning of new environments without physical hardware changes.

It also enhances reliability and security. By isolating workloads in separate VMs, it prevents system failures or security breaches in one environment from affecting others. And supports fast backup and recovery through snapshotting and migration features.

AWS DynamoDB is a fully managed NoSQL database service designed for applications requiring fast and consistent performance at scale. It automatically handles data replication, scaling, and backup, making it ideal for real-time analytics and e-commerce systems. Developers benefit from its low-latency reads and writes, flexible schema design, and seamless integration with other AWS services. However, DynamoDB's main limitations include higher costs at large scales, limited query flexibility compared to relational databases, and potential complexity when optimizing partition keys for performance.

The screenshot shows the AWS DynamoDB console interface. On the left, there's a navigation sidebar with options like Dashboard, Tables, Explore items, PartiQL editor, Backups, Exports to S3, Imports from S3, Integrations, Reserved capacity, and Settings. Under 'Explore items', 'TestUsers' is listed, and 'TestUsers1' is selected. The main area shows a search bar with 'Find tables' and a dropdown for 'Scan' or 'Query'. Below that, 'Select a table or index' is set to 'Table - TestUsers1' and 'Select attribute projection' is set to 'All attributes'. A section for 'Filters - optional' is present with a note to 'Add a filter to get started'. At the bottom are 'Run' and 'Reset' buttons. To the right, a table titled 'Table: TestUsers1 - Items returned (2)' is displayed. It shows two rows: one for 2021 with title 'Jujutsu Kaisen' and info 'Gojo is sealed.', and one for 2020 with title 'Demon Slayer' and info 'Nothing happens at all.'. There are 'Actions' and 'Create item' buttons above the table, and navigation controls (back, forward, search) at the top right.

year (Number)	title (String)	info
2021	Jujutsu Kaisen	Gojo is sealed.
2020	Demon Slayer	Nothing happens at all.

Amazon Simple Queue Service is a managed message queuing service that enables decoupled communication between distributed systems. It allows asynchronous message exchange, improving reliability and fault tolerance in event-driven architectures. SQS supports both standard and FIFO (First-In-First-Out) queues, ensuring flexibility for various use cases such as background task processing or microservice coordination. Its limitations include eventual consistency in standard queues, potential message duplication, and added latency compared to direct API calls. Despite these trade-offs, SQS remains a powerful tool for building scalable and resilient cloud applications.

The screenshot shows the AWS SQS console interface. At the top, there's a navigation bar with the AWS logo, a search bar, and account information (Account ID: 3023-4006-3392, United States (Ohio)). Below the navigation is a breadcrumb trail: Amazon SQS > Queues > test-quick-queue > Send and receive messages.

The main area has two sections:

- Amazon SQS** (left sidebar): Includes a "Queues" section and a message detail view for a received message.
- test-quick-queue** (right sidebar): Shows message attributes, a message content editor, and a message list table.

Message Detail View (top right):

- Header: Received message: 69dca163-eda1-4b31-b1c1-51ddf550cb45
- Tabs: Body (selected), Attributes, Details
- Content: Hello from AWS SQS! This is a test message.
- Buttons: Clear content, Send message, Done

Message List Table (bottom right):

- Header: Messages (1)
- Search bar: Search messages
- Table columns: ID, Sent, Size, Receive count
- Data: 1 message listed
- Buttons: View details, Delete, Poll for messages, Stop polling
- Footer: Polling progress: 0 receives/second